# Session 4: Data Transformation II — Pen-and-Paper Pair Exercise

**PSY 410 | Data Science for Psychology**

*No laptop today? No problem. This handout lets you practice the same skills on paper. Work with a partner who has a laptop and compare your work at the end.*

---

## The data: `flights`

Here are 12 rows from the `flights` dataset, showing just `carrier` and `dep_delay`:

| carrier | dep_delay |
|---|---|
| AA | 8 |
| AA | -2 |
| AA | 14 |
| DL | -3 |
| DL | 5 |
| DL | -1 |
| F9 | 25 |
| F9 | 38 |
| F9 | -4 |
| UA | 11 |
| UA | -6 |
| UA | 22 |

**Key:** Negative delays mean the flight departed early. Delays are in minutes.

---

## The task (same as the slide exercise)

1. Calculate the **average departure delay** for each carrier
2. Which airline has the **worst** average delay?
3. **Bonus:** Also calculate the number of flights per carrier. Does the worst airline just have fewer flights?

### Your pen-and-paper version

**Step 1: Group the data.** Draw a line between each group of rows by carrier. (They're already sorted for you in the table above.)

**Step 2: Summarize each group.** Calculate the mean `dep_delay` for each carrier. Show your work:

| carrier | dep_delay values | sum | n | mean (sum / n) |
|---|---|---|---|---|
| AA | | | | |
| DL | | | | |
| F9 | | | | |
| UA | | | | |

**Which carrier has the worst (highest) average delay?** _____

**Does that carrier just have fewer flights?** _____

**Step 3: Write the code.** Fill in the blanks to produce this summary:

```
flights |>
  ____________(carrier) |>
  ____________(
    avg_delay = _______(dep_delay, na.rm = TRUE),
    n_flights = ___()
  ) |>
  arrange(_______(avg_delay))
```

**Step 4: Think about it.** Why do we need `na.rm = TRUE` inside `mean()`?

Your answer: _____

---

### Check your work

Compare your summary table and code with your partner's screen.

**Hand calculations from the sample:**

| carrier | dep_delay values | sum | n | mean |
|---|---|---|---|---|
| AA | 8, -2, 14 | 20 | 3 | 6.67 |
| DL | -3, 5, -1 | 1 | 3 | 0.33 |
| F9 | 25, 38, -4 | 59 | 3 | 19.67 |
| UA | 11, -6, 22 | 27 | 3 | 9.00 |

**Worst carrier in the sample:** F9 (19.67 min average delay)

**From the full dataset:** F9 is also the worst overall (20.2 min avg delay), but with only 685 flights. EV is close behind (20.0 min) with 54,173 flights — so high delay isn't always just a small-sample problem.

**Expected code:**

```
flights |>
  group_by(carrier) |>
  summarize(
    avg_delay = mean(dep_delay, na.rm = TRUE),
    n_flights = n()
  ) |>
  arrange(desc(avg_delay))
```

**Why `na.rm = TRUE`?** Some flights have missing (`NA`) departure delay values. If we don't tell `mean()` to remove them, the whole average becomes `NA`. The `na.rm = TRUE` argument says "ignore the missing values and compute the mean from the rest."