

# Data Tidying

PSY 410: Data Science for Psychology

Dr. Sara Weston

2026-04-13

## What is tidy data?

Can you plot this?

| participant | depression_pre | depression_post | anxiety_pre | anxiety_post |
|-------------|----------------|-----------------|-------------|--------------|
| P1          | 18             | 12              | 24          | 20           |
| P2          | 22             | 14              | 19          | 15           |
| P3          | 15             | 11              | 28          | 22           |

Not easily. The columns mix variables (depression vs anxiety) with time points (pre vs post). ggplot doesn't know what to put on each axis.

By the end of today, you'll be able to reshape this into a form that works — and understand *why* it needs reshaping.

## The tidyverse philosophy

“Tidy datasets are all alike, but every messy dataset is messy in its own way.” — Hadley Wickham

The tools we've learned (ggplot2, dplyr) expect data in a specific format: **tidy data**.

## The three rules of tidy data

1. Each **variable** is a column
2. Each **observation** is a row
3. Each **value** is a cell

...  
Simple in theory, surprisingly complex in practice.

## Tidy data visualized

Table 2: Is this tidy?

| participant | pre_test | post_test |
|-------------|----------|-----------|
| P1          | 45       | 62        |
| P2          | 52       | 58        |
| P3          | 48       | 71        |

...  
**No!** Time (pre/post) is a variable, but it's spread across columns.

## The tidy version

Table 3: Now it's tidy!

| participant | time | score |
|-------------|------|-------|
| P1          | pre  | 45    |
| P1          | post | 62    |
| P2          | pre  | 52    |
| P2          | post | 58    |
| P3          | pre  | 48    |
| P3          | post | 71    |

Now: participant, time, and score are all columns.

## Why does it matter?

Tidy data works with tidyverse tools:

```
# Easy to visualize
ggplot(data, aes(x = time, y = score, color = participant)) +
  geom_point()

# Easy to analyze
data |>
  group_by(time) |>
  summarize(mean = mean(score))
```

## Wide vs. Long

### Wide format

- Variables spread across columns
- One row per subject
- Humans like to read this

### Long format

- Variables in a single column
- Multiple rows per subject
- R likes to work with this

Most real data needs reshaping.

## Common untidy patterns

### Pattern 1: Column headers are values

```
wide_scores <- tibble(
  student = c("Alice", "Bob", "Carol"),
  fall_2024 = c(85, 78, 92),
  spring_2025 = c(88, 82, 95),
  fall_2025 = c(91, 85, 94)
)
wide_scores
```

```
# A tibble: 3 x 4
  student fall_2024 spring_2025 fall_2025
  <chr>     <dbl>      <dbl>      <dbl>
1 Alice       85        88        91
2 Bob         78        82        85
3 Carol       92        95        94
```

The semester names are **values**, not variable names.

### Pattern 2: Multiple variables in one column

```
messy_data <- tibble(
  id = 1:3,
  age_sex = c("25_M", "32_F", "28_F")
)
messy_data
```

```
# A tibble: 3 x 2
  id age_sex
  <int> <chr>
1 1 25_M
2 2 32_F
3 3 28_F
```

Age and sex are crammed into one column.

### Pattern 3: Variables in rows and columns

```
weather <- tibble(
  id = c("MX001", "MX001", "MX002", "MX002"),
  year = c(2020, 2020, 2020, 2020),
  month = c(1, 2, 1, 2),
  element = c("tmax", "tmax", "tmin", "tmin"),
  value = c(85, 87, 32, 35)
)
weather
```

```
# A tibble: 4 x 5
  id      year month element value
  <chr>  <dbl> <dbl> <chr>   <dbl>
1 MX001    2020     1 tmax     85
2 MX001    2020     2 tmax     87
3 MX002    2020     1 tmin     32
4 MX002    2020     2 tmin     35
```

element contains variable names (tmax, tmin).

## Psychology-specific patterns

Surveys often look like:

```
survey_wide <- tibble(
  participant = 1:3,
  bdi_1 = c(2, 1, 3),
  bdi_2 = c(1, 0, 2),
  bdi_3 = c(3, 2, 2),
  bdi_4 = c(2, 1, 1)
)
survey_wide
```

```
# A tibble: 3 x 5
  participant bdi_1 bdi_2 bdi_3 bdi_4
  <int> <dbl> <dbl> <dbl> <dbl>
1 1         2     1     3     2
2 2         1     0     2     1
3 3         3     3     2     1
```

Each item is a column — wide format.

## pivot\_longer()

### The most common tidying operation

pivot\_longer() takes wide data and makes it long:

```

wide_scores |>
  pivot_longer(
    cols = fall_2024:fall_2025, # Which columns to pivot
    names_to = "semester",      # New column for old column names
    values_to = "score"        # New column for values
  )

```

```

# A tibble: 9 x 3
  student semester   score
  <chr>   <chr>     <dbl>
1 Alice    fall_2024     85
2 Alice    spring_2025   88
3 Alice    fall_2025     91
4 Bob      fall_2024     78
5 Bob      spring_2025   82
6 Bob      fall_2025     85
7 Carol   fall_2024     92
8 Carol   spring_2025   95
9 Carol   fall_2025     94

```

## Breaking it down

```

pivot_longer(
  cols = ...,           # Columns to reshape (use select helpers!)
  names_to = "...",    # Name for the new "names" column
  values_to = "..."    # Name for the new "values" column
)

```

## Selecting columns to pivot

Use any of the `select()` helpers:

```

# By name
pivot_longer(cols = c(fall_2024, spring_2025, fall_2025))

# By range
pivot_longer(cols = fall_2024:fall_2025)

# By pattern

```

```

pivot_longer(cols = starts_with("fall"))
pivot_longer(cols = contains("202"))

# Everything except
pivot_longer(cols = -student)

```

## Psychology example: Survey items

```

survey_wide |>
  pivot_longer(
    cols = starts_with("bdi"),
    names_to = "item",
    values_to = "response"
  )

```

```

# A tibble: 12 x 3
  participant item  response
  <int> <chr>   <dbl>
1       1 bdi_1     2
2       1 bdi_2     1
3       1 bdi_3     3
4       1 bdi_4     2
5       2 bdi_1     1
6       2 bdi_2     0
7       2 bdi_3     2
8       2 bdi_4     1
9       3 bdi_1     3
10      3 bdi_2     2
11      3 bdi_3     2
12      3 bdi_4     1

```

Now each response is its own row!

## Extracting information from names

What if column names contain useful info? We want to extract the time point (t1, t2, t3).

```
# Scores at different time points
experiment_wide <- tibble(
  id = 1:3,
  score_t1 = c(100, 95, 110),
  score_t2 = c(105, 100, 115),
  score_t3 = c(108, 102, 120)
)
experiment_wide
```

```
# A tibble: 3 x 4
  id score_t1 score_t2 score_t3
  <int>    <dbl>    <dbl>    <dbl>
1     1      100      105      108
2     2      95       100      102
3     3     110      115      120
```

### **names\_prefix argument**

```
experiment_wide |>
  pivot_longer(
    cols = starts_with("score"),
    names_to = "time",
    names_prefix = "score_", # Remove this prefix from names
    values_to = "score"
  )
```

```
# A tibble: 9 x 3
  id time  score
  <int> <chr> <dbl>
1     1 t1     100
2     1 t2     105
3     1 t3     108
4     2 t1     95
5     2 t2     100
6     2 t3     102
7     3 t1    110
8     3 t2    115
9     3 t3    120
```

## **names\_pattern argument**

For more complex parsing:

```
# Column names like "bdi_1", "anxiety_1", etc.  
multi_scale <- tibble(  
  id = 1:2,  
  bdi_1 = c(2, 1), bdi_2 = c(1, 2),  
  anxiety_1 = c(3, 2), anxiety_2 = c(2, 3)  
)  
multi_scale
```

```
# A tibble: 2 x 5  
  id   bdi_1  bdi_2  anxiety_1  anxiety_2  
  <int> <dbl> <dbl>     <dbl>      <dbl>  
1     1       2       1         3         2  
2     2       1       2         2         3
```

## **names\_pattern argument**

For more complex parsing:

```
multi_scale |>  
  pivot_longer(  
    cols = -id,  
    names_to = c("scale", "item"),  
    names_pattern = "(.+)_(.+)", # Regex: anythingAnything  
    values_to = "response"  
)
```

```
# A tibble: 8 x 4  
  id   scale  item  response  
  <int> <chr> <chr>    <dbl>  
1     1   bdi    1        2  
2     1   bdi    2        1  
3     1 anxiety  1        3  
4     1 anxiety  2        2  
5     2   bdi    1        1  
6     2   bdi    2        2  
7     2 anxiety  1        2  
8     2 anxiety  2        3
```

## Pivoting to calculate scale scores

```
# Calculate BDI total from long format
survey_wide |>
  pivot_longer(
    cols = starts_with("bdi"),
    names_to = "item",
    values_to = "response"
  ) |>
  group_by(participant) |>
  summarize(bdi_total = sum(response))
```

```
# A tibble: 3 x 2
  participant bdi_total
  <int>      <dbl>
1 1           8
2 2           4
3 3           8
```

## Pair coding break

**Your turn: 10 minutes**

We defined `survey_wide` earlier in this deck:

```
survey_wide
```

```
# A tibble: 3 x 5
  participant bdi_1 bdi_2 bdi_3 bdi_4
  <int> <dbl> <dbl> <dbl> <dbl>
1 1           2     1     3     2
2 2           1     0     2     1
3 3           3     2     2     1
```

With a partner, figure out: Which participant has the highest mean BDI score?

### 💡 Tip

You'll need `pivot_longer()` followed by `group_by()` + `summarize()`. The column names start with "bdi" — that's a hint.

---

## Before we move on

Upload your code to Canvas for participation credit. Paste what you have into today's in-class submission — it doesn't need to work perfectly.

## **pivot\_wider()**

### The opposite operation

Sometimes you need to go from long to wide:

```
long_data <- tibble(  
  participant = rep(1:3, each = 2),  
  time = rep(c("pre", "post"), 3),  
  score = c(45, 62, 52, 58, 48, 71)  
)  
long_data
```

```
# A tibble: 6 x 3  
  participant time   score  
  <int> <chr> <dbl>  
1           1 pre     45  
2           1 post    62  
3           2 pre     52  
4           2 post    58  
5           3 pre     48  
6           3 post    71
```

### **pivot\_wider() syntax**

```

long_data |>
  pivot_wider(
    names_from = time,    # Column to get new column names from
    values_from = score   # Column to get values from
  )

```

```

# A tibble: 3 x 3
  participant  pre  post
  <int> <dbl> <dbl>
1         1     45     62
2         2     52     58
3         3     48     71

```

## When to use pivot\_wider()

- Creating summary tables for reports
- Some analyses need wide format
- Merging data that was collected differently
- Human-readable output

## Multiple value columns

```

# Long data with multiple measures
long_multi <- tibble(
  id = rep(1:2, each = 2),
  time = rep(c("pre", "post"), 2),
  score = c(45, 62, 52, 58),
  rt = c(500, 480, 520, 490)
)
long_multi

```

```

# A tibble: 4 x 4
  id time  score     rt
  <int> <chr> <dbl> <dbl>
1     1 pre      45    500
2     1 post     62    480
3     2 pre      52    520
4     2 post     58    490

```

## Multiple value columns

```
long_multi |>
  pivot_wider(
    names_from = time,
    values_from = c(score, rt) # Multiple columns!
  )
```

```
# A tibble: 2 x 5
  id score_pre score_post rt_pre rt_post
  <int>     <dbl>      <dbl>   <dbl>    <dbl>
1     1        45        62     500     480
2     2        52        58     520     490
```

## Separating and uniting

**separate\_wider\_delim()**

Split one column into multiple:

```
tibble(
  id = 1:3,
  age_sex = c("25_M", "32_F", "28_F")) |>
  separate_wider_delim(
    cols = age_sex,
    delim = "_",
    names = c("age", "sex")
  )
```

```
# A tibble: 3 x 3
  id age   sex
  <int> <chr> <chr>
1     1 25    M
2     2 32    F
3     3 28    F
```

## **separate\_wider\_regex()**

For complex patterns:

```
tibble(
  code = c("A123", "B456", "C789")
) |>
  separate_wider_regex(
    cols = code,
    patterns = c(letter = "[A-Z]", number = "[0-9]+")
)
```

```
# A tibble: 3 x 2
  letter number
  <chr>  <chr>
1 A       123
2 B       456
3 C       789
```

## **unite()**

The opposite — combine columns:

```
tibble(
  year = c(2024, 2024, 2025),
  month = c(1, 6, 1),
  day = c(15, 20, 10)
) |>
  unite(
    col = "date",           # New column name
    year, month, day,      # Columns to combine
    sep = "-"               # Separator
)
```

```
# A tibble: 3 x 1
  date
  <chr>
1 2024-1-15
2 2024-6-20
3 2025-1-10
```

## Real-world examples

### Example 1: Repeated measures experiment

```
# Data as you might receive it from SPSS
wide_rm <- tibble(
  subject = 1:4,
  cond_a_time1 = c(450, 520, 480, 510),
  cond_a_time2 = c(420, 490, 460, 480),
  cond_b_time1 = c(480, 540, 500, 530),
  cond_b_time2 = c(440, 510, 470, 500)
)
wide_rm
```

```
# A tibble: 4 x 5
  subject cond_a_time1 cond_a_time2 cond_b_time1 cond_b_time2
    <int>      <dbl>      <dbl>      <dbl>      <dbl>
1       1        450        420        480        440
2       2        520        490        540        510
3       3        480        460        500        470
4       4        510        480        530        500
```

### Tidying repeated measures

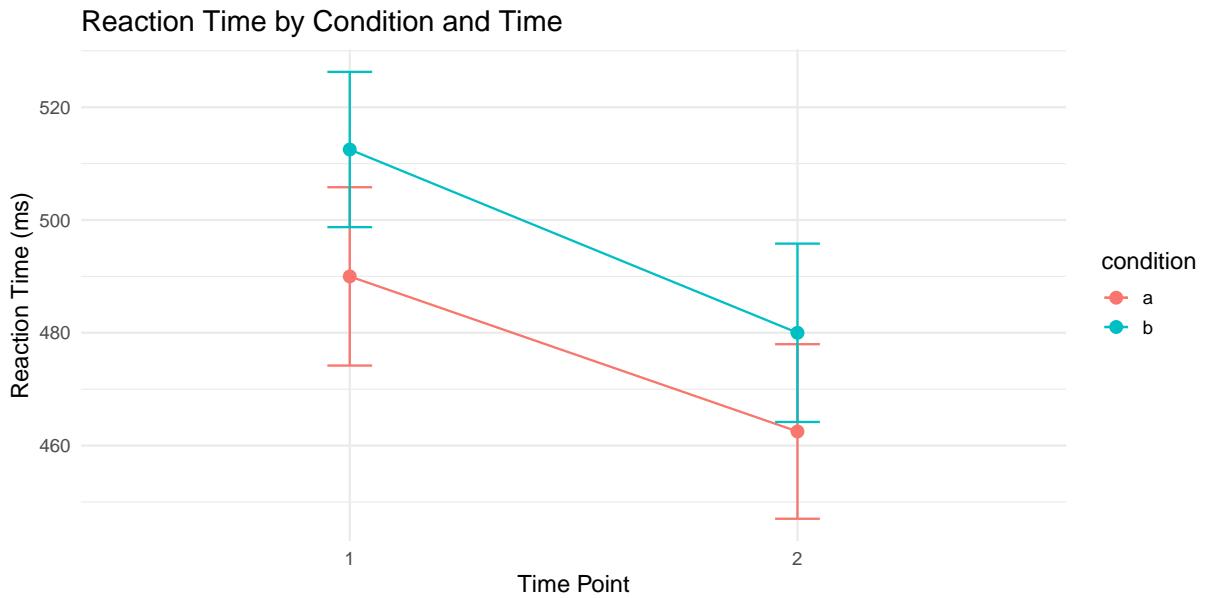
```
tidy_rm <- wide_rm |>
  pivot_longer(
    cols = -subject,
    names_to = c("condition", "time"),
    names_pattern = "cond_(.+)_(.+)",
    values_to = "rt"
)
tidy_rm
```

```
# A tibble: 16 x 4
  subject condition time     rt
    <int>   <chr>    <chr>  <dbl>
1       1     a        1      450
2       1     a        2      420
```

|    |     |   |     |
|----|-----|---|-----|
| 3  | 1 b | 1 | 480 |
| 4  | 1 b | 2 | 440 |
| 5  | 2 a | 1 | 520 |
| 6  | 2 a | 2 | 490 |
| 7  | 2 b | 1 | 540 |
| 8  | 2 b | 2 | 510 |
| 9  | 3 a | 1 | 480 |
| 10 | 3 a | 2 | 460 |
| 11 | 3 b | 1 | 500 |
| 12 | 3 b | 2 | 470 |
| 13 | 4 a | 1 | 510 |
| 14 | 4 a | 2 | 480 |
| 15 | 4 b | 1 | 530 |
| 16 | 4 b | 2 | 500 |

**Now we can analyze it!**

```
tidy_rm |>
  ggplot(aes(x = time, y = rt, color = condition, group = condition)) +
  stat_summary(fun = mean, geom = "point", size = 3) +
  stat_summary(fun = mean, geom = "line") +
  stat_summary(fun.data = mean_se, geom = "errorbar", width = 0.1) +
  labs(
    title = "Reaction Time by Condition and Time",
    x = "Time Point",
    y = "Reaction Time (ms)"
  ) +
  theme_minimal(base_size = 14)
```



### Example 2: Questionnaire with subscales

```
# Raw questionnaire data
quest <- tibble(
  pid = 1:3,
  anx_1 = c(3, 2, 4), anx_2 = c(2, 3, 3), anx_3 = c(4, 2, 5),
  dep_1 = c(2, 1, 3), dep_2 = c(3, 2, 4), dep_3 = c(2, 1, 3)
)

quest
```

```
# A tibble: 3 x 7
  pid  anx_1  anx_2  anx_3  dep_1  dep_2  dep_3
  <int> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1     1      3      2      4      2      3      2
2     2      2      3      2      1      2      1
3     3      4      3      5      3      4      3
```

### Example 2: Questionnaire with subscales

```

# Tidy and calculate subscales
quest |>
  pivot_longer(
    cols = -pid,
    names_to = c("scale", "item"),
    names_pattern = "(.+)_(.+)",
    values_to = "response"
  ) |>
  group_by(pid, scale) |>
  summarize(subscale_mean = mean(response), .groups = "drop") |>
  pivot_wider(names_from = scale, values_from = subscale_mean)

```

### Example 2: Questionnaire with subscales

```

# A tibble: 3 x 3
  pid   anx   dep
  <int> <dbl> <dbl>
1     1     3   2.33
2     2   2.33  1.33
3     3     4   3.33

```

### Example 3: Multilevel/nested data

```

# Students nested in classrooms
students <- tibble(
  classroom = rep(c("A", "B"), each = 3),
  student = 1:6,
  pretest = c(70, 75, 72, 68, 71, 69),
  posttest = c(80, 82, 78, 75, 79, 77)
)

students

```

```

# A tibble: 6 x 4
  classroom student pretest posttest
  <chr>      <int>   <dbl>    <dbl>
1 A            1       70      80
2 A            2       75      82
3 A            3       72      78

```

|   |   |   |    |    |
|---|---|---|----|----|
| 4 | B | 4 | 68 | 75 |
| 5 | B | 5 | 71 | 79 |
| 6 | B | 6 | 69 | 77 |

### Example 3: Multilevel/nested data

```
# Tidy for analysis
students |>
  pivot_longer(
    cols = c(pretest, posttest),
    names_to = "time",
    values_to = "score"
  )

# A tibble: 12 x 4
  classroom student time     score
  <chr>      <int> <chr>   <dbl>
1 A             1 pretest  70
2 A             1 posttest 80
3 A             2 pretest  75
4 A             2 posttest 82
5 A             3 pretest  72
6 A             3 posttest 78
7 B             4 pretest  68
8 B             4 posttest 75
9 B             5 pretest  71
10 B            5 posttest 79
11 B            6 pretest  69
12 B            6 posttest 77
```

**pivot\_longer() trips everyone up the first time**

### Pitfall 1: Forgetting what's a variable

Ask yourself: What are my variables?

- Participant ID? Variable
- Time point? Variable (not separate columns!)
- Score? Variable
- Item number? Depends on your analysis

## Pitfall 2: Over-pivoting

Not everything needs to be long:

```
# Maybe this is fine as-is?  
tibble(  
  id = 1:3,  
  age = c(25, 32, 28),  
  gender = c("M", "F", "F"),  
  score = c(85, 92, 88)  
)
```

Age, gender, and score are **different variables** — keep them as columns.

## General tidying strategy

1. **Identify** the variables (what are you measuring?)
2. **Look** at your current structure (what's a row? column?)
3. **Determine** what operations you need
4. **Test** with a small subset first
5. **Verify** you haven't lost data

## Get a head start

### Try it yourself

We created `wide_rm` earlier — a repeated measures dataset:

On your own:

1. Pivot it to long format, extracting **condition** and **time** from the column names
2. Calculate **mean RT by condition and time**
3. Sketch (on paper or in ggplot) what you'd expect the plot to look like

This is very close to what Assignment 3 will ask you to do.

# Wrapping up

## The `tidyverse` toolkit

| Function                    | What it does    |
|-----------------------------|-----------------|
| <code>pivot_longer()</code> | Wide → Long     |
| <code>pivot_wider()</code>  | Long → Wide     |
| <code>separate_*</code> ()  | Split columns   |
| <code>unite()</code>        | Combine columns |

## The tidy data mantra

1. Each **variable** is a column
2. Each **observation** is a row
3. Each **value** is a cell

When in doubt, ask: “What would make this easiest to plot/analyze?”

## Before next class

### Read:

- [R4DS Ch 7: Data import](#)
- [R4DS Ch 20: Spreadsheets](#)

### Practice:

- Reshape a dataset you’ve worked with
- Try tidying some messy example data
- Practice `pivot_longer()` — it’s the most common

## Key takeaways

1. **Tidy data** has a specific structure that works with `tidyverse`
2. `pivot_longer()` is your most-used tidying function
3. `pivot_wider()` is useful for tables and some analyses
4. **Think about your variables** before reshaping
5. **Column names contain information** — extract it with `names_pattern`

## **The one thing to remember**

If your data isn't tidy, your analysis can't start. `pivot_longer()` is the bridge.

Next time: Data Import