

Data Transformation I

PSY 410: Data Science for Psychology

Dr. Sara Weston

2026-04-06

Data transformation

A question you can't answer yet

You collected survey data from 500 participants. You only want women over 25 who passed the attention check.

...

How do you get to *just those rows*?

...

That's what `filter()` does. And it's just the beginning — today we learn four verbs that turn raw data into exactly what you need.

The dplyr verbs

Verb	What it does
<code>filter()</code>	Pick rows by their values
<code>arrange()</code>	Reorder rows
<code>select()</code>	Pick columns by name
<code>mutate()</code>	Create new columns
<code>summarize()</code>	Collapse to a summary

Today: `filter()`, `arrange()`, `select()`, `mutate()`

Our dataset: flights

```
# All 336,776 flights departing NYC in 2013
glimpse(flights)
```

Rows: 336,776
Columns: 19

```
$ year           <int> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2~  
$ month          <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~  
$ day            <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~  
$ dep_time        <int> 517, 533, 542, 544, 554, 554, 555, 557, 557, 558, 558, ~  
$ sched_dep_time <int> 515, 529, 540, 545, 600, 558, 600, 600, 600, 600, ~  
$ dep_delay       <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2, -2, -2, -2, -~  
1~  
$ arr_time        <int> 830, 850, 923, 1004, 812, 740, 913, 709, 838, 753, 849, ~  
$ sched_arr_time  <int> 819, 830, 850, 1022, 837, 728, 854, 723, 846, 745, 851, ~  
$ arr_delay        <dbl> 11, 20, 33, -18, -25, 12, 19, -14, -8, 8, -2, -3, 7, -~  
1~  
$ carrier         <chr> "UA", "UA", "AA", "B6", "DL", "UA", "B6", "EV", "B6", "~  
$ flight          <int> 1545, 1714, 1141, 725, 461, 1696, 507, 5708, 79, 301, 4~  
$ tailnum         <chr> "N14228", "N24211", "N619AA", "N804JB", "N668DN", "N394~  
$ origin          <chr> "EWR", "LGA", "JFK", "JFK", "LGA", "EWR", "EWR", "LGA", ~  
$ dest            <chr> "IAH", "IAH", "MIA", "BQN", "ATL", "ORD", "FLL", "IAD", ~  
$ air_time         <dbl> 227, 227, 160, 183, 116, 150, 158, 53, 140, 138, 149, 1~  
$ distance         <dbl> 1400, 1416, 1089, 1576, 762, 719, 1065, 229, 944, 733, ~  
$ hour             <dbl> 5, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 6, 5, 6, 6, 6~  
$ minute           <dbl> 15, 29, 40, 45, 0, 58, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 59, 0~  
$ time_hour        <dttm> 2013-01-01 05:00:00, 2013-01-01 05:00:00, 2013-01-~  
01 0~
```

Understanding the data

- `year`, `month`, `day` — departure date
- `dep_time`, `arr_time` — actual times (HHMM format)
- `dep_delay`, `arr_delay` — delays in minutes (negative = early)
- `carrier` — airline code
- `origin`, `dest` — airport codes
- `air_time`, `distance` — in minutes and miles

filter()

filter() picks rows

```
# All flights on January 1st
filter(flights, month == 1, day == 1)

# A tibble: 842 x 19
   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
   <int> <int> <int>     <int>          <int>      <dbl>    <int>          <int>
1  2013     1     1      517            515        2     830          819
2  2013     1     1      533            529        4     850          830
3  2013     1     1      542            540        2     923          850
4  2013     1     1      544            545       -1    1004         1022
5  2013     1     1      554            600       -6     812          837
6  2013     1     1      554            558       -4     740          728
7  2013     1     1      555            600       -5     913          854
8  2013     1     1      557            600       -3     709          723
9  2013     1     1      557            600       -3     838          846
10 2013     1     1      558            600       -2     753          745
# i 832 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>
```

Comparison operators

Operator	Meaning
==	equal to
!=	not equal to
<, >	less than, greater than
<=, >=	less/greater than or equal



Use == for comparison, not !=

Filter examples

```
# Flights to Los Angeles
filter(flights, dest == "LAX")

# Flights with arrival delay over 2 hours
filter(flights, arr_delay > 120)

# Flights by United Airlines
filter(flights, carrier == "UA")
```

Multiple conditions

Conditions separated by , are combined with AND:

```
# January 1st flights (both must be true)
filter(flights, month == 1, day == 1)
```

Equivalent to:

```
filter(flights, month == 1 & day == 1)
```

Logical operators

Operator	Meaning
&	AND (both true)
	OR (either true)
!	NOT (negation)

Using OR

```
# Flights in November OR December
filter(flights, month == 11 | month == 12)
```

```

# A tibble: 55,403 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>        <int>     <dbl>    <int>        <int>
1 2013    11     1      5        2359       6     352        345
2 2013    11     1     35        2250      105     123       2356
3 2013    11     1    455        500      -5     641       651
4 2013    11     1    539        545      -6     856       827
5 2013    11     1    542        545      -3     831       855
6 2013    11     1    549        600     -11     912       923
7 2013    11     1    550        600     -10     705       659
8 2013    11     1    554        600      -6     659       701
9 2013    11     1    554        600      -6     826       827
10 2013   11     1    554        600      -6     749       751
# i 55,393 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

```

A useful shortcut: %in%

```

# Same as above, but cleaner
filter(flights, month %in% c(11, 12))

```

```

# A tibble: 55,403 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>        <int>     <dbl>    <int>        <int>
1 2013    11     1      5        2359       6     352        345
2 2013    11     1     35        2250      105     123       2356
3 2013    11     1    455        500      -5     641       651
4 2013    11     1    539        545      -6     856       827
5 2013    11     1    542        545      -3     831       855
6 2013    11     1    549        600     -11     912       923
7 2013    11     1    550        600     -10     705       659
8 2013    11     1    554        600      -6     659       701
9 2013    11     1    554        600      -6     826       827
10 2013   11     1    554        600      -6     749       751
# i 55,393 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

```

`%in%` checks if the value is in the vector.

filter() with psychology data

Imagine a survey dataset:

```
# Exclude participants who failed attention check
filter(survey, attention_check == "correct")

# Keep only complete responses
filter(survey, !is.na(total_score))

# Adults only
filter(survey, age >= 18)

# Specific conditions
filter(survey, condition %in% c("treatment", "control"))
```

Missing values: NA

`NA` means “not available” — a missing value.

```
x <- c(1, 2, NA, 4)
x > 2
```



```
[1] FALSE FALSE     NA  TRUE
```

Any operation with `NA` returns `NA` (it’s unknown!).

Checking for NA

Use `is.na()` to check for missing values:

```
x <- c(1, 2, NA, 4)
is.na(x)
```



```
[1] FALSE FALSE  TRUE FALSE
```

```
# Keep rows where dep_delay is NOT missing  
filter(flights, !is.na(dep_delay))
```

Pair coding break

Your turn: 10 minutes

With a partner, using the `flights` dataset:

1. Find all **United Airlines** ("UA") flights
2. that were **more than 2 hours late** arriving
3. and were flying **to Los Angeles** ("LAX")

How many flights match? Which origin airport had the most?



Tip

You'll need `filter()` with multiple conditions. Think about which operators you need.

Before we move on

Upload your code to Canvas for participation credit. Paste what you have into today's in-class submission — it doesn't need to work perfectly.

arrange()

arrange() reorders rows

```
# Sort by departure delay (smallest first)  
arrange(flights, dep_delay)
```

```

# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>          <int>     <dbl>    <int>        <int>
1 2013     12     7    2040            2123      -43       40        2352
2 2013      2     3    2022            2055      -33      2240        2338
3 2013     11    10    1408            1440      -32      1549        1559
4 2013      1    11    1900            1930      -30      2233        2243
5 2013      1    29    1703            1730      -27      1947        1957
6 2013      8     9     729             755      -26      1002        955
7 2013     10    23    1907            1932      -25      2143        2143
8 2013      3    30    2030            2055      -25      2213        2250
9 2013      3     2    1431            1455      -24      1601        1631
10 2013      5     5     934             958      -24      1225        1309
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

```

Descending order

```

# Most delayed flights first
arrange(flights, desc(dep_delay))

```

```

# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>          <int>     <dbl>    <int>        <int>
1 2013     1     9     641             900      1301      1242        1530
2 2013      6    15    1432            1935      1137      1607        2120
3 2013     10    10    1121            1635      1126      1239        1810
4 2013     20     9    1139            1845      1014      1457        2210
5 2013     22     7     845             1600      1005      1044        1815
6 2013     10     4    1100            1900      960       1342        2211
7 2013     17     3    2321             810      911       135         1020
8 2013     27     6     959             1900      899       1236        2226
9 2013     22     7    2257             759      898       121         1026
10 2013      5    12    756             1700      896      1058        2020
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

```

Multiple sort columns

```
# Sort by month, then by day, then by departure time
arrange(flights, month, day, dep_time)

# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>     <int>        <int>    <dbl> <int>        <int>
1 2013     1     1      517         515      2     830        819
2 2013     1     1      533         529      4     850        830
3 2013     1     1      542         540      2     923        850
4 2013     1     1      544         545     -1    1004       1022
5 2013     1     1      554         600     -6     812        837
6 2013     1     1      554         558     -4     740        728
7 2013     1     1      555         600     -5     913        854
8 2013     1     1      557         600     -3     709        723
9 2013     1     1      557         600     -3     838        846
10 2013    1     1      558         600     -2     753        745
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>
```

arrange() use cases

```
# Psychology examples:
# Find participants with highest scores
arrange(data, desc(total_score))

# Sort by condition, then by participant ID
arrange(data, condition, participant_id)

# Find earliest responses
arrange(data, response_time)
```

Quick mental model

You've now seen three verbs. Each answers a different question:

Verb	Question it answers
<code>filter()</code>	Which rows do I keep?
<code>arrange()</code>	What order should they be in?
<code>select()</code>	Which columns do I need?
<code>mutate()</code>	What new columns should I create?

We've done the first two. Now let's pick our columns.

select()

select() picks columns

```
# Just these three columns
select(flights, year, month, day)
```

```
# A tibble: 336,776 x 3
  year month   day
  <int> <int> <int>
1 2013     1     1
2 2013     1     1
3 2013     1     1
4 2013     1     1
5 2013     1     1
6 2013     1     1
7 2013     1     1
8 2013     1     1
9 2013     1     1
10 2013    1     1
# i 336,766 more rows
```

Select a range

```
# All columns from year to day
select(flights, year:day)
```

```
# A tibble: 336,776 x 3
  year month   day
  <int> <int> <int>
1 2013     1     1
2 2013     1     1
3 2013     1     1
4 2013     1     1
5 2013     1     1
6 2013     1     1
7 2013     1     1
8 2013     1     1
9 2013     1     1
10 2013    1     1
# i 336,766 more rows
```

Select helpers

Helper	What it does
<code>starts_with("x")</code>	Columns starting with “x”
<code>ends_with("x")</code>	Columns ending with “x”
<code>contains("x")</code>	Columns containing “x”
<code>everything()</code>	All remaining columns

Using select helpers

```
# All delay-related columns
select(flights, contains("delay"))
```

```
# A tibble: 336,776 x 2
  dep_delay arr_delay
  <dbl>      <dbl>
1       2        11
2       4        20
3       2        33
4      -1       -18
5      -6       -25
6      -4        12
7      -5        19
```

```

8       -3      -14
9       -3      -8
10      -2       8
# i 336,766 more rows

```

Using select helpers

```

# All time columns
select(flights, ends_with("time"))

# A tibble: 336,776 x 5
  dep_time sched_dep_time arr_time sched_arr_time air_time
  <int>        <int>     <int>        <int>     <dbl>
1     517         515     830         819      227
2     533         529     850         830      227
3     542         540     923         850      160
4     544         545    1004        1022      183
5     554         600     812         837      116
6     554         558     740         728      150
7     555         600     913         854      158
8     557         600     709         723      53
9     557         600     838         846      140
10    558         600     753         745      138
# i 336,766 more rows

```

Reorder columns

```

# Move air_time and distance to the front
select(flights, air_time, distance, everything())

```

```

# A tibble: 336,776 x 19
  air_time distance year month   day dep_time sched_dep_time dep_delay
  <dbl>     <dbl> <int> <int> <int> <int>        <int>     <dbl>
1     227     1400  2013    1     1     517         515      2
2     227     1416  2013    1     1     533         529      4
3     160     1089  2013    1     1     542         540      2
4     183     1576  2013    1     1     544         545     -1
5     116      762  2013    1     1     554         600     -6

```

```

6      150      719 2013     1     1      554      558     -4
7      158     1065 2013     1     1      555      600     -5
8       53      229 2013     1     1      557      600     -3
9      140      944 2013     1     1      557      600     -3
10     138      733 2013     1     1      558      600     -2
# i 336,766 more rows
# i 11 more variables: arr_time <int>, sched_arr_time <int>, arr_delay <dbl>,
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

```

Exclude columns

Use - to remove columns:

```
# Remove year column
select(flights, -year)
```

```

# A tibble: 336,776 x 18
  month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int>    <int>          <int>    <dbl>    <int>          <int>
1     1     1      517            515      2     830          819
2     1     1      533            529      4     850          830
3     1     1      542            540      2     923          850
4     1     1      544            545     -1    1004         1022
5     1     1      554            600     -6     812          837
6     1     1      554            558     -4     740          728
7     1     1      555            600     -5     913          854
8     1     1      557            600     -3     709          723
9     1     1      557            600     -3     838          846
10    1     1      558            600     -2     753          745
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

```

Exclude columns

```
# Remove multiple columns
select(flights, -year, -month, -day)
select(flights, -(year:day)) # same thing
```

select() in psychology

```
# Select demographic columns
select(survey, starts_with("demo_"))

# Select all BDI items
select(survey, contains("bdi"))

# Remove identifying information
select(survey, -name, -email, -ip_address)

# Reorder for analysis
select(survey, participant_id, condition, starts_with("outcome"))
```

The pattern so far

Notice: every verb has the same shape.

```
# verb(data, what_youWant)
filter(flights, month == 1)
arrange(flights, dep_delay)
select(flights, year, month, day)
```

Data goes first. Then you describe what you want. This consistency is by design.

mutate()

mutate() creates new columns

```
# Calculate total delay
mutate(flights, total_delay = dep_delay + arr_delay)
```

```
# A tibble: 336,776 x 20
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>          <int>     <dbl>    <int>          <int>
1 2013     1     1      517            515       2     830           819
2 2013     1     1      533            529       4     850           830
```

```

3 2013 1 1 542 540 2 923 850
4 2013 1 1 544 545 -1 1004 1022
5 2013 1 1 554 600 -6 812 837
6 2013 1 1 554 558 -4 740 728
7 2013 1 1 555 600 -5 913 854
8 2013 1 1 557 600 -3 709 723
9 2013 1 1 557 600 -3 838 846
10 2013 1 1 558 600 -2 753 745
# i 336,766 more rows
# i 12 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>, total_delay <dbl>

```

Multiple new columns

```

mutate(flights,
  # Create new variables
  total_delay = dep_delay + arr_delay,
  speed = distance / air_time * 60, # mph
  # Can reference columns you just created!
  delay_per_mile = total_delay / distance
)

# A tibble: 336,776 x 22
   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
   <int> <int> <int>    <int>        <int>     <dbl>    <int>        <int>
1 2013     1     1      517          515       2     830        819
2 2013     1     1      533          529       4     850        830
3 2013     1     1      542          540       2     923        850
4 2013     1     1      544          545      -1    1004       1022
5 2013     1     1      554          600      -6     812        837
6 2013     1     1      554          558      -4     740        728
7 2013     1     1      555          600      -5     913        854
8 2013     1     1      557          600      -3     709        723
9 2013     1     1      557          600      -3     838        846
10 2013    1     1      558          600      -2     753        745
# i 336,766 more rows
# i 14 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>, total_delay <dbl>, speed <dbl>,
# delay_per_mile <dbl>

```

mutate() keeps all columns

`mutate()` adds new columns while keeping existing ones.

Use `transmute()` to *only* keep new columns:

```
transmute(flights,
          total_delay = dep_delay + arr_delay,
          speed = distance / air_time * 60
)
```

```
# A tibble: 336,776 x 2
  total_delay speed
  <dbl>     <dbl>
1       13     370.
2       24     374.
3       35     408.
4      -19     517.
5      -31     394.
6        8     288.
7       14     404.
8      -17     259.
9      -11     405.
10      6     319.
# i 336,766 more rows
```

Useful functions in mutate()

Arithmetic: `+`, `-`, `*`, `/`, `^`, `%%` (modulo), `%/%` (integer division)

Logs: `log()`, `log2()`, `log10()`

Offsets: `lead()`, `lag()` (for time series)

Cumulative: `cumsum()`, `cummean()`, `cummax()`

Ranking: `min_rank()`, `dense_rank()`, `row_number()`

Common transformations

```

# Z-score (standardize)
mutate(data, score_z = (score - mean(score)) / sd(score))

# Log transform (for skewed data)
mutate(data, rt_log = log(reaction_time))

# Create categories from continuous
mutate(data, age_group = case_when(
  age < 30 ~ "young",
  age < 60 ~ "middle",
  TRUE ~ "older"
))

```

mutate() for psychology

```

# Calculate scale scores (mean of items)
mutate(survey,
       bdi_total = (bdi_1 + bdi_2 + bdi_3 + bdi_4) / 4,
       # Or rowwise if you have many items:
       anxiety = rowMeans(select(., anx_1:anx_20), na.rm = TRUE)
)

# Create dummy codes
mutate(survey,
       female = if_else(gender == "female", 1, 0),
       treatment = if_else(condition == "treatment", 1, 0)
)

# Reverse code items
mutate(survey,
       item_5r = 8 - item_5 # For 1-7 scale
)

```

The pipe makes your code read like a sentence

The problem with nested functions

What if we want to:

1. Filter to January flights
2. Select departure time and delay
3. Arrange by delay

Nested approach:

```
arrange(
  select(
    filter(flights, month == 1),
    dep_time, dep_delay
  ),
  dep_delay
)
```

This is hard to read!

Intermediate objects approach

```
# Save each step
jan_flights <- filter(flights, month == 1)
jan_selected <- select(jan_flights, dep_time, dep_delay)
jan_arranged <- arrange(jan_selected, dep_delay)
```

Works, but clutters your environment with objects.

The pipe: |>

The pipe takes the result of one function and passes it as the first argument to the next:

```
# Same result, much cleaner!
flights |>
  filter(month == 1) |>
  select(dep_time, dep_delay) |>
  arrange(dep_delay)
```

```
# A tibble: 27,004 x 2
  dep_time dep_delay
     <int>     <dbl>
1      1900      -30
```

```
2      1703      -27
3      1354      -22
4      2137      -22
5       704      -21
6     2050      -20
7     2134      -20
8     2050      -20
9     2140      -19
10    1947      -18
# i 26,994 more rows
```

Reading piped code

Read `|>` as “then”:

```
flights |>                      # Start with flights, THEN
  filter(month == 1) |>          # filter to January, THEN
  select(dep_time, dep_delay) |> # select these columns, THEN
  arrange(dep_delay)            # arrange by delay
```

Keyboard shortcut

The pipe is so common, there's a shortcut:

- **Windows/Linux:** Ctrl + Shift + M
- **Mac:** Cmd + Shift + M

...



Tip

In RStudio settings, make sure “Use native pipe operator” is enabled (Tools → Global Options → Code)

Note: `|>` vs `%>%`

You'll see both:

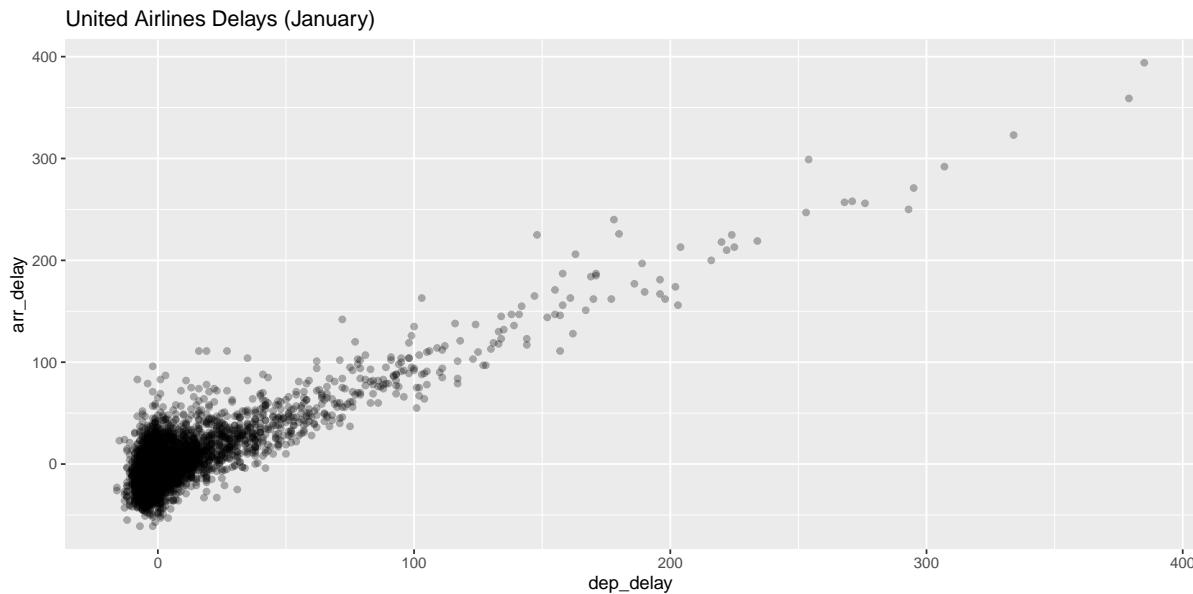
- `|>` — the **native** pipe (built into R 4.1+)
- `%>%` — the **magrittr** pipe (older, from tidyverse)

They work almost identically. We'll use `|>` since it's now standard.

Piping into ggplot

The pipe works beautifully with ggplot:

```
flights |>
  filter(month == 1, carrier == "UA") |>
  ggplot(aes(x = dep_delay, y = arr_delay)) +
  geom_point(alpha = 0.3) +
  labs(title = "United Airlines Delays (January)")
```



Putting it together

A complete workflow

```
# Which airlines have the longest delays in summer?
flights |>
  filter(month %in% c(6, 7, 8)) |>                      # Summer months
  filter(!is.na(arr_delay)) |>                            # Remove NAs
  mutate(delay_hours = arr_delay / 60) |>                 # Convert to hours
  select(carrier, delay_hours) |>                          # Keep relevant columns
  arrange(desc(delay_hours))                                # Longest delays first
```

```
# A tibble: 84,124 x 2
  carrier delay_hours
  <chr>     <dbl>
1 MQ         18.8
2 MQ         16.5
3 DL         14.9
4 DL         14.2
5 AA         13.4
6 DL          13
7 DL         12.8
8 VX          11.3
9 AA          10.8
10 VX         10.5
# i 84,114 more rows
```

Get a head start

Your turn!

Using the `flights` dataset:

1. Filter to American Airlines (“AA”) flights to Los Angeles (“LAX”)
2. Create a new variable `speed` (`distance / air_time * 60`)
3. Select carrier, origin, dest, and speed
4. Arrange by speed (highest first)
5. What’s the fastest AA flight to LAX?

Wrapping up

The dplyr workflow

```
data |>
  filter(<conditions>) |>      # Pick rows
  select(<columns>) |>        # Pick columns
  mutate(<new vars>) |>       # Create columns
  arrange(<order>)           # Sort rows
```

The pipe (`|>`) connects these verbs into a readable workflow.

Before next class

Read:

- [R4DS Ch 3: Data transformation](#) (section 3.5)
- [R4DS Ch 4: Workflow: code style](#)

Practice:

- Transform `flights` in different ways
- Create new variables with `mutate()`
- Build multi-step pipelines

Key takeaways

1. `filter()` picks rows by condition
2. `arrange()` reorders rows (use `desc()` for descending)
3. `select()` picks columns (helpers like `contains()` are useful)
4. `mutate()` creates new columns
5. The pipe `|>` makes code readable and elegant

The one thing to remember

The pipe turns a wall of nested code into a sentence you can read aloud.

Next time: `group_by()` and `summarize()`