

Exploratory Data Analysis: Covariation

PSY 410: Data Science for Psychology

Dr. Sara Weston

2026-05-04

From variation to covariation

Psychology is about relationships

Last time, you explored how individual variables behave — distributions, outliers, missing data.

But psychology is about *relationships*:

- Does **treatment** predict **depression**?
- Does **age** relate to **reaction time**?
- Does **anxiety** co-occur with **insomnia**?

Today we learn to **see** those relationships in data — before testing them statistically.

Categorical + Continuous

Example: Mental health by treatment group

```
# Simulated therapy outcome data
therapy_data <- tibble(
  condition = rep(c("Control", "CBT", "Mindfulness"), each = 50),
  depression_post = c(
    rnorm(50, mean = 18, sd = 5), # Control
    rnorm(50, mean = 12, sd = 5), # CBT
    rnorm(50, mean = 14, sd = 5)  # Mindfulness
  )
)
```

Example: Mental health by treatment group

```
glimpse(therapy_data)
```

Rows: 150

Columns: 2

\$ condition <chr> "Control", "Control", "Control", "Control",
"Control",...


\$ depression_post <dbl> 26.559783, 18.316721, 13.498250, 28.065713,
16.112715,...

Boxplots: The classic choice


```
ggplot(therapy_data,  
      aes(x = condition, y = depression_post, fill = condition)) +  
  geom_boxplot(alpha = 0.7, show.legend = FALSE) +  
  scale_fill_manual(values = c(  
    "Control" = "#0072B2", "CBT" = "#E69F00", "Mindfulness" = "#009E73")) +  
  labs(  
    title = "CBT shows the lowest depression scores",  
    x = NULL,  
    y = "Depression score (BDI-II)"  
  ) +  
  theme_minimal(base_size = 14) +  
  theme(panel.grid.major.x = element_blank())
```

Boxplots: The classic choice

CBT shows the lowest depression scores



What boxplots show



- **Line in middle:** median
- **Box:** 25th–75th percentile (IQR)
- **Whiskers:** $1.5 \times \text{IQR}$
- **Dots:** outliers

...

Great for comparison — but they hide the actual data.

The problem with boxplots

Boxplots summarize, but they hide important information:

- **The actual distribution shape** (is it bimodal? skewed?)
- **Individual data points** (how many observations are there?)
- **The raw data** (where do specific values fall?)

We can do better.

Raincloud plots

The modern psych visualization

Raincloud plots combine three elements:


1. **Density curve** (distribution shape)
2. **Boxplot** (summary stats)
3. **Jittered points** (individual data)

They're increasingly popular in psychology publications because they show **everything**.

Use the **gg****rain** package. ([More details here.](#))

Building a raincloud

Post-Treatment Depression by Condition
Diamond = mean. Density = distribution. Box = IQR. Dots = individual scores.



```

library(gggrain)
therapy_data |>
  ggplot(aes(
    x = condition,
    y = depression_post,
    fill = condition,
    color = condition)) +
  # geom_rain creates all parts of your raincloud
  geom_rain(
    alpha = .6,
    # change just the boxplot part
    boxplot.args = list(color = "black")) +
  # The mean
  stat_summary(fun = mean, geom = "point", shape = 18, size = 5, color = "black") +
  scale_fill_manual(
    values = c("Cognitive Behavioral Therapy" = "#00728F", "Mindfulness-Based Cognitive Therapy" = "#E69F00", "Transcendental Meditation" = "#000000", "Control Group" = "#000000")
  )

```

When to use which

Boxplot

- Median, IQR, outliers
- Fast to read
- Best for **quick EDA** and large datasets where individual points would clutter

Raincloud

- Density + summary + raw data
- Shows shape, sample size, and outliers all at once
- Best for **publications and presentations**

Pair coding break

Your turn: Build a raincloud

Using `therapy_data`, build a **raincloud plot** of depression scores by condition.

1. Load `ggrain`
2. Map `x = condition`, `y = depression_post`, and `fill = condition`
3. Add `geom_rain()`
4. Give it an informative title and labeled axes

Time: 10 minutes

Before we move on

 Submit your code on [Canvas](#) for participation credit. Paste what you have — it doesn't need to work perfectly.

Categorical + Categorical

Example: Diagnosis by referral source

```
# Simulated diagnostic data
diagnosis_data <- tibble(
  referral = sample(c("Self", "Physician", "School"), 200, replace = TRUE),
  diagnosis = sample(c("Depression", "Anxiety", "Both", "Other"),
                    200, replace = TRUE)
)
```


```
head(diagnosis_data)
```

```
# A tibble: 6 × 2
  referral diagnosis
  <chr>      <chr>
1 Self      Anxiety
2 Physician Depression
3 Self      Other
4 Self      Depression
5 Physician Both
6 School    Other
```

Option 1: Stacked bar chart

```
ggplot(diagnosis_data, aes(x = referral, fill = diagnosis)) +  
  geom_bar() +  
  labs(  
    title = "Diagnosis by referral source",  
    x = "Referral source",  
    y = "Count",  
    fill = "Diagnosis"  
  ) +  
  theme_minimal()
```


Option 1: Stacked bar chart



Option 2: Side-by-side bars

```
ggplot(diagnosis_data, aes(x = referral, fill = diagnosis)) +  
  geom_bar(position = "dodge") +  
  labs(  
    title = "Diagnosis by referral source",  
    x = "Referral source",  
    y = "Count",  
    fill = "Diagnosis"  
  ) +  
  theme_minimal()
```


Option 2: Side-by-side bars



Option 3: Proportions

```
ggplot(diagnosis_data, aes(x = referral, fill = diagnosis)) +  
  geom_bar(position = "fill") +  
  labs(  
    title = "Diagnosis distribution by referral source",  
    x = "Referral source",  
    y = "Proportion",  
    fill = "Diagnosis"  
  ) +  
  theme_minimal()
```

Option 3: Proportions



Continuous + Continuous

The classic: Scatterplots

```
# Simulated reaction time data
rt_data <- tibble(
  age = runif(100, 18, 70),
  reaction_time = 200 + age * 3 + rnorm(100, 0, 40)
)

glimpse(rt_data)
```

Rows: 100

Columns: 2


\$ age <dbl> 38.00245, 37.89727, 67.53066, 63.82036, 18.76507,
47.552...

\$ reaction_time <dbl> 255.5022, 408.5319, 362.0651, 427.6740, 268.7969,
326.06...

Basic scatterplot

```
ggplot(rt_data, aes(x = age, y = reaction_time)) +  
  geom_point() +  
  labs(  
    title = "Reaction time by age",  
    x = "Age (years)",  
    y = "Reaction time (ms)"  
  ) +  
  theme_minimal()
```


Basic scatterplot



Add a trend line

```
ggplot(rt_data, aes(x = age, y = reaction_time)) +  
  geom_point() +  
  geom_smooth(method = "lm", se = TRUE) +  
  labs(  
    title = "Reaction time increases with age",  
    x = "Age (years)",  
    y = "Reaction time (ms)"  
  ) +  
  theme_minimal()
```

Add a trend line



The problem: Overplotting

With lots of data, points overlap and hide the true density:


```
# Lots of data
big_rt_data <- tibble(
  age = runif(5000, 18, 70),
  reaction_time = 200 + age * 3 + rnorm(5000, 0, 40)
)
```

Overplotting problem demonstrated

```
ggplot(big_rt_data, aes(x = age, y = reaction_time)) +  
  geom_point() +  
  labs(title = "Hard to see where the data is dense") +  
  theme_minimal()
```

Overplotting problem demonstrated

Hard to see where the data is dense




Solution 1: Transparency (alpha)

```
ggplot(big_rt_data, aes(x = age, y = reaction_time)) +  
  geom_point(alpha = 0.1) +  
  labs(title = "Using alpha = 0.1 to show density") +  
  theme_minimal()
```

Solution 1: Transparency (alpha)


Using alpha = 0.1 to show density



Solution 2: geom_bin2d()

```
ggplot(big_rt_data, aes(x = age, y = reaction_time)) +  
  geom_bin2d() +  
  scale_fill_viridis_c() +  
  labs(  
    title = "2D bins show density",  
    fill = "Count"  
  ) +  
  theme_minimal()
```

Solution 2: geom_bin2d()



Correlation coefficient

A single number summary of the linear relationship:

```
cor(rt_data$age, rt_data$reaction_time)
```

```
[1] 0.7604132
```

Note

- **r = 1**: perfect positive relationship
- **r = 0**: no linear relationship
- **r = -1**: perfect negative relationship

But always plot your data first! (See: Anscombe's Quartet)

Patterns and models

What patterns tell us

When you see covariation, ask:

1. **Could it be coincidence?** (Maybe, especially with small samples)
2. **What's the mechanism?** (How are these variables related?)
3. **Is there a confound?** (Could a third variable explain both?)

Warning

Correlation $\neq$ Causation

Covariation suggests a relationship, but doesn't prove one variable *causes* the other.

Your class data: Covariation

Stress by chronotype

```
ggplot(class_survey, aes(x = chronotype, y = stress, fill = chronotype)) +  
  geom_boxplot(alpha = 0.7, show.legend = FALSE) +  
  labs(  
    title = "Are Night Owls More Stressed?",  
    x = "Chronotype",  
    y = "Current stress level (1–10)"  
  ) +  
  theme_minimal(base_size = 14)
```


Stress by chronotype



Coding excitement vs. coding anxiety

```
ggplot(class_survey, aes(x = coding_excited, y = coding_anxious)) +  
  geom_jitter(width = 0.2, height = 0.2, alpha = 0.6, size = 3) +  
  geom_smooth(method = "lm", se = FALSE, color = "steelblue") +  
  labs(  
    title = "Does Excitement Buffer Anxiety About Coding?",  
    x = "Excitement about learning to code (1–10)",  
    y = "Anxiety about learning to code (1–10)"  
  ) +  
  theme_minimal(base_size = 14)
```


Coding excitement vs. coding anxiety



Caffeine and sleep

```
ggplot(class_survey, aes(x = caffeine_per_day, y = sleep_hrs)) +  
  geom_jitter(width = 0.2, height = 0.2, alpha = 0.6, size = 3) +  
  geom_smooth(method = "lm", se = FALSE, color = "steelblue") +  
  labs(  
    title = "More Caffeine, Less Sleep?",  
    x = "Caffeinated drinks per day",  
    y = "Average hours of sleep per night"  
  ) +  
  theme_minimal(base_size = 14)
```


Caffeine and sleep



Wrapping up

Key takeaways

1. **Covariation = relationships** between variables
2. **Different plot types** for different variable combinations:
 - Categorical + continuous: boxplot, raincloud
 - Categorical + categorical: stacked, dodged, or proportional bars
 - Continuous + continuous: scatterplot (with alpha or 2D bins for big data)
3. **Rainclouds** show distribution, summary stats, and raw data all at once
4. **Watch for overplotting** — use alpha, jitter, or binning
5. **Always visualize first** before computing correlations
6. **Patterns suggest but don't prove** causation

Before next class

Read:

- R4DS Ch 12: Logical vectors
- R4DS Ch 13: Numbers

Do:

- Get started on Assignment 5 (due Sun May 10)
- Start thinking about how you'll compute scale scores — that's where we're headed Wednesday

The one thing to remember

Relationships hide in data. Your job is to make them visible — carefully, honestly.

See you Wednesday!

Get a head start

Assignment 5 preview

Assignment 5 asks you to do EDA on the **BFI** personality dataset. Two of the tasks line up directly with what we did today:

- **Task 2.2:** Scatterplot of two Extraversion items with a trend line
- **Task 2.3:** Boxplots of one Extraversion item by education level

If you have time tonight, take a first pass at either one — the code patterns are the same as today's class survey examples.